


Architecture overview

FDDT (Fraud Document Detection Tool) authenticates third-party documents — vendor invoices, school documents, quotations, travel claims and the payment evidence that supports them — *before* a payment or approval is made. A user submits a **case** containing one or more PDF documents. The system extracts their data, runs a battery of automated checks, turns the results into an explainable risk score, and hands the case to a human reviewer (L1, escalating to L2 when needed) who approves, rejects or escalates it. Every step is written to an append-only audit log.

Several client **companies** share one deployment; each company's users see only that company's data, enforced twice — in the application and by PostgreSQL Row-Level Security (multi-tenancy.md).

System diagram



Everything intelligent is a managed cloud API or a classical library. **No model is trained or hosted** (a design rule of the platform).

Tech stack

Layer	Technology	Used for
API	FastAPI 0.121, Pydantic 2, Uvicorn	REST endpoints, validation, OpenAPI
Auth	PyJWT (HS256), bcrypt	Email/password login, 8 h bearer tokens carrying <code>role</code> , <code>company_id</code> , <code>is_platform_admin</code>
Database	PostgreSQL 16, SQLAlchemy 2, Alembic, psycopg2	Persistence, migrations (14), Row-Level Security per company
Connection pooling	PgBouncer (transaction mode)	Multiplexes API + worker connections onto ≤ 60 server connections
Queue	Celery 5.4 + Redis 7	Three queues (extraction / vision / forensics), fair-share priorities, beat jobs
Rate limiting	Redis sliding windows (<code>services/rate_limiter.py</code> , <code>services/login_throttle.py</code>)	Global caps on Azure calls across every worker; sign-in throttling per email address and per IP
File store	Azure Blob Storage (<code>azure-storage-blob</code>)	Immutable originals, signature crops, report PDFs (company-prefixed paths)
OCR	Azure Document Intelligence <code>prebuilt-layout</code>	Text, tables, key-values and word geometry
LLM	Azure OpenAI (chat completions, strict JSON-schema output)	Classification, field extraction, issuer semantic match, visual review, signature detect/compare
PDF forensics	pikepdf, PyMuPDF	Upload validation, metadata, xref/incremental-update analysis, page rendering, report PDF
Image forensics	OpenCV (headless), Pillow, NumPy	ELA, copy-move (BRISK), ghost text (traces of erased text in converted scans), signature crop/pixel comparison
Duplicates	imagehash (pHash)	Near-duplicate page detection (within a company)
Fuzzy match	rapidfuzz	Issuer registry lookup, signer-name comparison
Workflow	<code>transitions</code>	Case state machine
Frontend	React 19, TypeScript 6, Vite 8, React Router 7	SPA, route-level code splitting
UI kit	Tailwind CSS 4, Radix primitives, shadcn-style <code>ui</code> components, lucide-react	Styling and base components
Data	TanStack Query 5	Server state, polling
Forms	React Hook Form + Zod	Form state and validation
PDF view	react-pdf (pdf.js)	Live viewer with overlays, signature-region drawing

Key architectural decisions

1. Modular monolith

One FastAPI app and the Celery workers share one codebase (`backend/app/`), one database and one settings object. Modules are separated by responsibility — `api/` (HTTP), `services/` (logic, mostly pure functions), `tasks/` (Celery wrappers that own DB sessions and audit events), `models/`, `schemas/`, `db/` (sessions and tenancy) — not by deployable unit. The same image runs as API, as each worker pool and as beat, differing only in the start command.

The convention that makes this work: **services contain the algorithm and are testable without a database or network; tasks wrap them with I/O**. For example `services/forensics/ela.py` `run_ela_check(pages)` is pure; `tasks/tampering_checks_task.py` downloads the file, renders it, calls the pure function and stores the result.

2. Azure OpenAI behind an `LLMService` abstraction

Classification, field extraction, issuer semantic matching, visual review and signature analysis all use Azure OpenAI. `services/llm_service.py` defines the `LLMService` interface and one implementation, `AzureOpenAILLMService`. Callers depend only on the interface, so switching to another LLM provider means adding one implementation class, not rewriting call sites. The same pattern guards `StorageService` (Azure Blob) and `OCRService` (Document Intelligence). All LLM calls request **strict JSON-schema output** (`response_format: json_schema`, `strict: true`), so every response is validated by a Pydantic model (`DocumentAnalysis`, `PageVisualAnalysis`, `SignatureComparisonResult`, ...). Vision calls require a vision-capable deployment (`gpt-4o`, `gpt-4.1`, `gpt-4.1-mini`, ...) and raise `LLMConfigurationError` otherwise.

3. Three queues, fair-share, globally rate-limited

Uploads return immediately. The upload endpoint validates the file, commits the document row and enqueues **one task per document per step** — six of them. Each task goes to the queue of the service it depends on: `extraction_queue` (Document Intelligence), `vision_queue` (Azure OpenAI) or `forensics_queue` (local CPU, prefork). Each queue has its own worker pool, so they scale independently. Dispatch is **fair-share per company**: a task's priority comes from how much work its company already has outstanding, so one company's 100-document batch can't starve another company's single upload. Azure calls pass through a **Redis-backed global rate limiter** (requests and tokens per minute, calibrated against real usage) shared by every worker. Tasks commit before every slow step (no transaction is held across a download, an Azure call or CPU work), and check results are written **idempotently** (one row per document + check), so a redelivered task overwrites rather than duplicates. The browser polls `GET /cases/{id}` (every 5 s while the pipeline is incomplete). Details: `processing-queues.md`.

3b. Event-driven case workflow

State changes go through `services/workflow_service.py`, which drives a `transitions` state machine and **emits** `WorkflowEvents` (`case_status_changed`, `case_approved`, `case_rejected`, `case_escalated`) to subscribers instead of calling side effects directly. The only subscriber today writes the audit log. A future external case-management integration would be one more `subscribe()` call plus the already-present nullable `cases.external_ref_id` column — no change to the state machine.

4. Immutable, validated originals

Only PDFs are accepted, and every upload is validated **before** it is stored: within the company's own size limit (per-company, set by a platform admin; 10 MB by default), not empty, a PDF by its header bytes, not corrupted, not password-protected (`pipeline/01-upload-intake.md`). A valid file is written to Blob Storage exactly once, at `companies/{company_id}/cases/{case_id}/documents/{document_id}_{sha256}{ext}`, with `overwrite=False`; `StorageService` exposes no update or delete. Its SHA-256 is stored on `documents.file_hash` so integrity can be re-verified later — the PDF report does exactly that and prints the result. Nothing downstream mutates the original: forensics render pages **in memory**; the signature crop and the report are *new* blobs (`.../signatures/...`, `.../reports/...`). The private container means the API never hands out the raw blob URL; it mints short-lived SAS URLs (15 minutes by default, 30 for OCR), and refuses to sign a path under another company's prefix.

5. One normalized bounding-box convention

Every located thing — ELA/copy-move regions, visual-review findings, extracted-field positions, signature/stamp regions, signature references — is stored as `{"page": 1-based int, "x", "y", "width", "height"}` where the four numbers are **fractions (0–1) of the page's width/height, origin top-left**. It is resolution-independent, so the same box works on a 200-DPI server render, a browser canvas, and the report. The frontend draws boxes live over the real

PDF (PdfOverlayViewer) and never saves an annotated image; only the exported report burns findings into its own rendered pages.

6. A transparent rules engine, not a model

Risk scoring is a weighted rules engine (services/risk_scoring_service.py). Each rule is a row in the company's risk_rules describing what to match in already-stored check results; there is no per-rule Python. Rules are immutable and versioned — editing inserts version + 1 — and every scoring run (case_risk_assessments) freezes the rendered reason text, the exact rule versions and the thresholds it used, so retuning never changes a historical score. New companies start from the platform rule template (risk_rule_templates). See pipeline/11-risk-scoring-engine.md.

7. Template-free extraction

One generic pipeline handles every document type: Azure Layout OCR → a single Azure OpenAI call that both classifies the document and extracts core fields plus arbitrary additional fields as {name, value} pairs. No per-type extractors exist, so new document types need no engineering.

8. Two-layer tenant isolation

Every tenant-owned table carries company_id . The application binds each request's (and each task's) database session to one company and filters every ORM query automatically; independently, PostgreSQL Row-Level Security refuses any row of another company even to SQL with no filter at all. No database role has BYPASSRLS; the platform path sees all rows through an explicit platform_all policy. A cross-company id is a 404. See multi-tenancy.md.

9. Explainability over accuracy claims

Weak signals are labelled as such everywhere in code and UI: ELA is a "signal, not a verdict"; signature/stamp detection is presence-and-placement only; signature comparison uses advisory wording ("visually consistent with reference on file", never "verified"); the visual review's AI question is "experimental". Weights reflect that.

Scope of the current build

What each capability area does today, and what is intentionally outside the current scope.

#	Area	Implementation	Where
1	Classification & extraction	Azure OpenAI via LLMService , one combined call per document	services/llm_service.py
2	Signature/stamp detection	Azure OpenAI vision model locates signature and stamp regions (presence/placement only)	services/signature_detection_service.py
3	AI-generated-content assessment	One question inside the visual-review vision call; experimental, one signal among several. There is no separate content-detection service	services/visual_inconsistency_service.py
4	Metadata forensics	pikepdf + PyMuPDF read the PDF's own object model	services/forensics/metadata_forensics.py
5	Accepted file types	PDF only , within the company's per-file limit (10 MB by default), validated at upload (type by header bytes, corruption, password). Every forensic check is PDF-based	services/upload_validation.py

#	Area	Implementation	Where
5a	Bulk upload	One zip (within the company's zip limit, 300 MB by default) with one folder per case. Ingested in the background case by case, each document validated like a single upload and queued as soon as its case exists. Partial success, live per-case summary	<code>services/bulk_upload_service.py</code>
5b	Upload limits	Per company (<code>companies.max_file_size_mb</code> / <code>max_zip_size_mb</code>), changed only by a platform admin, audited old → new, effective on the next upload; one source for every upload path	<code>services/upload_limits.py</code>
6	Error Level Analysis	Runs only on pages containing an embedded raster image; born-digital text pages return <code>not_applicable</code>	<code>services/forensics/ela.py</code>
7	Case workflow	Statuses in use: <code>submitted</code> , <code>pending_manual_review</code> , <code>approved</code> , <code>rejected</code> . Every scored case goes to the reviewer queue; nothing auto-approves. Escalation moves the case to the L2 reviewer tier (<code>assigned_tier</code>), not a status	<code>services/workflow_service.py</code>
8	Low-confidence fields	The per-field <code>uncertain</code> flag is stored and displayed to the reviewer; it does not change routing, because every case reaches manual review	<code>services/llm_service.py</code>
9	Multi-tenancy	Companies, <code>company_id</code> everywhere, application filter + PostgreSQL RLS, platform admin with audited read-only support access	<code>db/tenancy.py</code> , <code>api/tenant_access.py</code>
10	Billing & usage	Per-company counters (cases, documents, files, storage) with nightly reconciliation; platform-admin dashboard	<code>services/usage_service.py</code>
11	Processing queues	Three queues, fair-share per company, global Azure rate limiter, queue monitor, JSON queue metrics + <code>queue_alert</code> log line	<code>tasks/celery_app.py</code> , <code>tasks/fairshare.py</code>
12	Notifications	Not part of the current scope (no notifications table, no email sending)	—
13	SLA tracking	Not part of the current scope (no deadline column, no overdue checks)	—
14	Organisation-wide MIS reporting	Not part of the current scope. The <i>per-case</i> forensic report is fully built	<code>services/case_report_service.py</code>
15	Frontend libraries	Tables are hand-built per page; no charting library is used	<code>frontend/package.json</code>
16	Signature verification	Detection + reviewer-created references + in-case vision comparison + a classical pixel-identical-reuse check. Cross-case library comparison is not built (the <code>is_library</code> flag only stores data)	<code>services/signature_*</code>
17	Evidence export	Per-case PDF report with file hashes and annotated renders; no bundle of the original files	<code>services/case_report_service.py</code>
18	Arabic values	Extracted Arabic field values render with the default text direction on the case detail page; Arabic issuer names render right-to-left in the Issuer Registry settings page	<code>frontend/src/pages/</code>

#	Area	Implementation	Where
19	Document labels	<code>school_document</code> , <code>vendor_invoice</code> , <code>commercial_invoice</code> , <code>procurement_documentation</code> , <code>quotation</code> , <code>travel_invoice</code> , <code>payment_evidence</code> , <code>other</code> . Note the case type is <code>travel_reimbursement</code> while the document label is <code>travel_invoice</code>	<code>services/classification_service.py</code> , <code>models/case.py</code>
20	Cloud deployment	Not built — local development only. The production plan and costs are in <code>deployment/production-deployment.md</code>	—

Security posture

- **Two mandatory checks on every endpoint:** a role check (`require_company_role` / `require_reader` with a minimum rank — `user < reviewer_l1 < reviewer_l2` — or `require_platform_admin`) and a tenant check (the request's session is bound to the caller's company by `get_tenant_db`). The user is re-read from the database on every request; a token whose company or platform flag no longer matches, or whose company is suspended, is refused. Case actions additionally check the case's tier (`app/api/case_access.py` `can_act_on_case`).
- **Tenant isolation is enforced twice** — application-layer filtering and PostgreSQL Row-Level Security — and a cross-company id is a **404**. Platform admins belong to no company, are read-only on company data, and every one of their accesses writes a platform-only audit row that no company user can see.
- **Submitters are deliberately blinded:** `GET /cases` and `GET /cases/{id}` return a neutral "In review" flag and no risk tier/assessment to a `user` , and the case timeline (`GET /cases/{id}/audit-log`) omits the scoring event and strips result details (results, finding and match counts, score, tier, rules fired) from the automated check events — so a bad actor cannot learn the score. A case owned by someone else is a **404**, not **403**. *Note:* the per-document check cards on the case page (e.g. "Issuer verification: Flagged") are still shown to the submitter — see the change report R6 for this open decision.
- **Case-content routes are guarded** (`api/case_access.py`). Document upload and `GET .../documents/{id}/file-url` require that the caller is the case's **owner or a reviewer of its company** — anyone else gets a **404** (`load_visible_case`), and a rejected upload writes nothing to storage. The three `signatures` endpoints follow the same rule: the **uploader sets the reference signature during upload**, so the case owner may use them, as may any reviewer; a stranger gets a 404. Covered by `tests/test_case_authorization.py` . A submitter can still read their own case's signature comparison results (`GET .../signature-matches`), unlike the risk assessment, which is withheld.
- **Uploads are validated before storage** (PDF only, within the company's size limit, not corrupted, not password-protected), so a malformed file never reaches Blob Storage, a worker or an Azure API.
- Tokens are stored in `localStorage` and there is no server-side logout, refresh or password-reset endpoint. `JWT_SECRET_KEY` defaults to `changeme-in-.env` — production must set it (see `deployment/production-deployment.md`).