

# 10 — Duplicate & near-duplicate detection

**What it detects:** a page that looks the same as a page in a document seen before — a resubmitted invoice, the same quotation re-uploaded under a slightly different vendor name, a re-scan, a re-saved copy. It catches what a byte-level SHA-256 comparison ( `documents.file_hash` ) cannot: re-saving, re-compressing or trivially editing a file changes its bytes but not how it *looks*.

**Code:** `services/forensics/duplicate_check.py` ( `run_duplicate_check` ), run by `tasks/duplicate_check_task.py` ( `run_duplicate_check_task` ); model `models/document_page_hash.py` ( `document_page_hashes` ); stores `check_type = duplicate_detection`.

## Algorithm

1. Render each page ( `render_pdf_pages` , 200 DPI) and compute a **perceptual hash**: `imagehash.phash` (64-bit, `hash_size = 8` ) of the RGB page image → a 16-character hex string, one per page ( `compute_page_hashes` ).
2. `find_matches` loads **every stored page hash of every other document of the same company** (joined to filename, case id and case number) and, for each new page, finds the single **closest** prior page whose Hamming distance is  $\leq$  `DUPLICATE_HASH_HAMMING_THRESHOLD` (default **5** of 64 bits — deliberately conservative: near- identical resubmissions, not merely similar layouts such as two invoices from one template). One match per new page at most.
3. `store_page_hashes` writes this document's own hashes **regardless of outcome**, so future uploads can match against it — as an upsert, one row per (document, page), so a redelivered task can't duplicate them. (The hash is a side effect of the check; there is no separate step.)
4. `result` is `flag` if any page matched, else `pass`.

Match details (filename, case number) are captured **at check time**, not resolved later, so a finding stays readable even if the matched document's case is one the current viewer has never opened.

## Output shape (real)

```
{
  "result": "flag",
  "details": [{
    "finding": "near_duplicate_page",
    "severity": "high",
    "page": 1,
    "description": "Page 1 is a near-identical match (perceptual hash Hamming distance 0) to page 1 of \"Sample15_Arabic_Case_Evidence.pdf\" in case CASE-8D8C00E1 – a possible duplicate or near-duplicate resubmission.",
    "data": {
      "matched_document_id": "uuid", "matched_document_filename": "Sample15_Arabic_Case_Evidence.pdf",
      "matched_case_id": "uuid", "matched_case_number": "CASE-8D8C00E1",
      "matched_page": 1, "distance": 0
    }
  ]
}
```

Same `{finding, severity, description, page, data}` shape as ELA/copy-move (no `bounding_box` — the whole page is the match), so the shared findings-list UI renders it. In the UI the finding links to the matched case.

`document_page_hashes` : `id` , `company_id` , `document_id` (FK), `page_number` (1-based), `phash` ( `varchar(32)` , indexed), timestamps; unique on ( `document_id` , `page_number` ).

## Field-aware: resubmission or the same template

A page hash sees layout, so different invoices printed on one template all match. Once both documents' fields are extracted (the check runs in parallel with extraction, so this happens after field validation, or straight away when the fields are already there — `refine_duplicate_check`), each page match that is not the very same file is compared on the invoice on that page ( `extracted_fields.invoices` for a multi-invoice file): invoice number, date, amount and student / customer.

- all compared fields equal (at least two) → stays `near_duplicate_page`, high: the same document submitted again ( `data.verdict` = "resubmission");
- any of them different → `same_template_page`, **low** (shown, not scored): the same template with other values, listed with the differences ( `data.field_comparison` );
- fields not available yet → left as it is.

## Known limitations

- **It matches against everything the company ever stored**, including your own test uploads. Re-uploading the same sample file in a new case (or uploading the two halves of one case that share a page) **will flag**. This is correct behaviour, and the most common source of "unexpected" flags during testing.
- **Never across companies**. The same invoice submitted to two different client companies is not detected: matching across tenants would reveal another company's filenames and case numbers. This is a deliberate privacy choice.
- **Same-case matches count** (see the note at the top) — despite the rule's `cross_case` name.
- Full-scan comparison per page: cost grows linearly with the number of stored pages. There is no index (e.g. BK-tree/LSH); fine for thousands of pages, not millions.
- pHash is sensitive to layout, not content: two documents that differ only in a changed amount will hash almost identically (that is the intent, and why it is high severity), whereas a heavy crop or rotation can defeat it.
- No cleanup: hashes of documents whose cases are closed stay forever. Two uploads processed at exactly the same moment by parallel forensics processes may not see each other.
- PDF only (as are uploads).

## Risk rules fed

| rule_id                                                          | Finding                          | Weight | Severity |
|------------------------------------------------------------------|----------------------------------|--------|----------|
| <code>duplicate.cross_case_match</code> (v2: resubmissions only) | <code>near_duplicate_page</code> | 40     | high     |

This is the heaviest single seeded rule: with the default thresholds (medium  $\geq 30$ , high  $\geq 60$ ) one duplicate page alone makes a case **medium** risk; with any other signal it approaches **high**.