

Multi-tenancy

FDDT serves several client companies from one deployment. Each company's users see **only** that company's data; the platform operator's own team (platform admins) sits outside every company.

Roles

Role	Company	Can
<code>user</code>	exactly one	submit cases/documents, track own cases
<code>reviewer_1</code>	exactly one	review the company's queue; approve / reject / escalate (L1 tier)
<code>reviewer_2</code>	exactly one	everything L1 can, plus escalated (L2) cases, plus the company's own Settings > Issuer Registry and Settings > Risk Rules
<code>platform_admin</code>	none	create companies; create every user and assign company + role; edit the platform risk-rule template new companies start from; billing & usage dashboard; processing-queue monitor; read-only, audited support access to any company's cases, audit log and settings (and may edit a company's issuer registry/risk rules on its behalf, also audited)

- Company roles are ranked (`has_rank` / `hasRank`); `platform_admin` is deliberately *not* on that ladder, so it never inherits "reviewer" powers — it cannot approve, reject, escalate, upload, create cases or generate reports.
- There is no company-level user management: **Settings > Users** is platform-admin only.
- The former single-tenant `admin` role was migrated to `platform_admin`.

Data model

- `companies (id, name, is_active, created_at, updated_at)`. `is_active = false` suspends a company: its users can't sign in and existing tokens stop working on the next request.
- `company_id` (FK, indexed) on every tenant-owned table: `users`, `cases`, `documents`, `document_checks`, `document_page_hashes`, `cross_document_findings`, `case_actions`, `case_risk_assessments`, `risk_scores`, `risk_rules`, `risk_settings`, `issuer_registry`, `signature_references`, `signature_matches`, `case_reports`, `audit_log`.
 - NOT NULL everywhere except `users` (NULL = platform admin; a CHECK constraint enforces `role = 'platform_admin' ⇔ company_id IS NULL`) and `audit_log` (NULL = platform-level event such as *company created, template edited, usage reconciled, platform-admin access*). *User created / updated* for a company user is written into **that company's** log; the company sees the actor as "Platform team" (the platform admin's own user row is invisible to a company session).
 - Extracted fields are a column on `documents` and duplicate-detection hashes live in `document_page_hashes` — both covered by those tables' `company_id`.
- `risk_rules`, `risk_settings` (tier thresholds) and `issuer_registry` are **per company**. Rule uniqueness is `(company_id, rule_id, version)`. The issuer registry starts empty.

- `risk_rule_templates` (platform-level, **no** `company_id`) holds the default rule set. Platform admins edit it under **Platform > Rule templates**. When a company is created, every **active** template row is copied into that company's own `risk_rules` (version 1, `company_id` stamped) plus default thresholds (`seed_risk_rules(db, company_id)`). From then on the copy is the company's own: editing the template never changes an existing company — only companies created afterwards (the same "no retroactive change" principle as rule versioning). The migration seeded the template from the Default Company's current rules (38 rules — its 42 rows were 38 rules plus 4 older versions). Only `fddt_platform` has access to the table.
- `company_usage_stats` (`company_id`, `usage_date`, `cases_created`, `documents_uploaded`, `files_stored`, `storage_bytes`) — see [Billing & usage](#).

Tenant isolation — two independent layers

1. Application layer (`app/db/tenancy.py` , `app/api/auth.py` , `app/api/tenant_access.py`)

- The JWT carries `role` , `company_id` (null for platform admins) and `is_platform_admin` . Every request re-reads the user from the database; if the token's company or platform flag no longer matches (user moved to another company, role changed across the platform boundary) the token is refused (401), and a suspended company's users are refused too.
- Every protected endpoint applies **two mandatory checks**: the role check (`require_company_role` , `require_reader` , `require_platform_admin`) **and** the tenant check — the request's data session is bound to the caller's company by the `get_tenant_db` dependency, so a route cannot forget it.
- Every session is in one of three states:
 - **unbound** — refuses to query or write any tenant-owned table (raises `TenantContextError`);
 - **company** — every ORM SELECT/UPDATE/DELETE touching a tenant model gets `company_id = <bound company>` added automatically (`with_loader_criteria` , including relationship loads and `session.get()`); every INSERT is stamped with that company and a row of another company is refused at flush; objects of other companies are evicted from the identity map on binding;
 - **platform** — the explicit RLS-bypassing path (below), no automatic filter.
- On top of that backstop, service functions that read tenant data take `company_id` explicitly and filter by it themselves (`load_current_rules` , `get_risk_settings` , `match_issuer` , `run_duplicate_check` , `score_case` , `pipeline_status` , `latest_assessment` , `get_case_flags` , ...).
- An id that belongs to another company is a **404** — never a silent empty result and never a 403 that would confirm the id exists.
- Celery tasks are enqueued as `task.delay(document_id, company_id)` and open a session bound to that company (`app/tasks/tenant.py`), so workers are confined exactly like requests.
- Duplicate detection compares a document only against **its own company's** history (matching across companies would leak another tenant's filenames and case numbers).

2. Database layer — PostgreSQL Row-Level Security

The migration `a7c1e9d3f5b2_multi_tenancy` creates two login roles and the policies:

Role	Attributes	Used for
<code>fddt_app</code> (<code>DATABASE_APP_USER</code>)	NOSUPERUSER NOBYPASSRLS	every company-scoped request and pipeline task
<code>fddt_platform</code> (<code>DATABASE_PLATFORM_USER</code>)	NOSUPERUSER NOBYPASSRLS — full access through an explicit <code>platform_all</code> policy on every table	sign-in/token identity lookup, platform-admin screens, usage reconciliation, resolving a legacy task's company

Role	Attributes	Used for
owner (DATABASE_URL)	owns the tables; needs CREATEROLE, not superuser	Alembic migrations and <code>seed.py</code> only

No role has BYPASSRLS, and no superuser is needed anywhere. Platform access is policy-based: `CREATE POLICY platform_all ON <table> FOR ALL TO fddt_platform USING (true) WITH CHECK (true)` on each RLS table — creating a policy only needs table ownership, so this works unchanged on managed PostgreSQL (Azure Flexible Server). Verified by running the **entire** migration chain as a NOSUPERUSER/CREATEROLE owner on a fresh database and re-running the RLS suite there. (PostgreSQL 16 requires superuser even to *mention* `SUPERUSER / BYPASSRLS` in `ALTER ROLE`, so those attributes are set on `CREATE ROLE` only and the migration asserts afterwards that the roles really are unprivileged.)

- RLS is enabled on every tenant-owned table (plus `companies` and `company_usage_stats`) with a `tenant_isolation` policy for `fddt_app`: `USING (company_id = fddt_current_company_id()) WITH CHECK (same)`, where `fddt_current_company_id()` reads the transaction-local setting `app.current_company_id`.
- The app sets that setting at the start of **every** transaction of a company-bound session (`after_begin` hook, `set_config(..., is_local => true)`), so a pooled connection never carries one request's company into the next. No setting $\Rightarrow$ no rows.
- A bug that forgets the company filter therefore still cannot read, update, delete or insert another company's rows — the database refuses. `tests/test_rls_postgres.py` proves this with raw SQL that has **no** company filter at all.
- Grants are least-privilege and also make the append-only/immutable tables read+insert-only **in the database** for both app roles: `audit_log`, `risk_rules` (versioned), `case_risk_assessments`, `case_reports`. `users` is SELECT-only for `fddt_app` (user management is platform-only).
- The platform engine is reachable only through a session explicitly bound to platform mode (`get_system_db`, `system_session()`); the default request session is unbound and then company-bound.

Adding a tenant-owned table later: give it `TenantScopedMixin`, and in its migration `ENABLE ROW LEVEL SECURITY`, create the `tenant_isolation` policy (`enable_tenant_rls()` in the multi-tenancy migration shows how) and grant the two app roles explicitly. There are deliberately no default privileges, so a new table is unreachable by the app until that is done.

Platform-admin support access (audited)

A platform admin can open any company's case, case list, audit log or settings. The route resolves which company the resource belongs to through the platform session, binds the request to **that** company (so the rest of the request is confined by both layers), and writes a `platform_admin_access` audit row, e.g.:

Platform admin `ops@example.com` viewed Company B / Case CASE-1A2B3C4D

These rows are platform-only. They are written through the platform session with `company_id = NULL` (the target company is in `event_data.company_id`), so:

- RLS hides them from every company session — a company query cannot return them even without a filter;
- every company-facing view (Audit History, its event-type filter, a case's timeline, the case report) also excludes the event type explicitly;
- company users have no indication that platform access happened.

Platform admins see every row — in the platform-level log, and in a company's log view (which shows the company's own events plus the platform-access rows about it). **Every access is recorded individually:**

`PLATFORM_ACCESS_AUDIT_DEDUP_SECONDS` defaults to `0` (a value > 0 would collapse identical reads).

Rows written before this change were moved to `company_id = NULL` by migration `c9e3a1b5d7f4` (no row removed, no content changed).

Platform admins have no cross-company case list: the case queue takes one company at a time.

Blob storage layout

```
companies/{company_id}/cases/{case_id}/documents/{document_id}_{sha256}{ext}
companies/{company_id}/cases/{case_id}/reports/{report_id}.pdf
companies/{company_id}/cases/{case_id}/signatures/{crop_id}.png
```

`ensure_company_blob()` refuses to sign a URL under another company's prefix. Files written before the migration keep their original paths (`{case_id}/...`, `reports/...`, `signatures/...`) — originals are immutable and are not moved; their database rows are company-scoped like everything else.

Billing & usage

Platform > Billing & usage shows one row per company (never a mixed list of cases): cases created, documents uploaded, files stored (documents + generated reports), storage added in the period, and each company's current all-time storage and file count, for this month / last month / last 30 days / this year / all time / a custom UTC date range.

- Counters live in `company_usage_stats`, one row per company per UTC day, bumped with an atomic `INSERT ... ON CONFLICT DO UPDATE SET n = n + delta` **in the same transaction** that creates the case, document or report (`app/services/usage_service.py`). The dashboard sums a few small rows — never a COUNT/SUM over the case tables.
- Storage is the exact byte size recorded when each file was written (`documents.file_size_bytes`, `case_reports.file_size_bytes`). Signature crops are small derived images and are not counted.
- `reconcile_usage_stats` (Celery beat, nightly at `USAGE_RECONCILIATION_HOUR_UTC`, default 02:00 UTC) recomputes every (company, day) from the source tables and corrects any drift, writing a platform-level `usage_stats_reconciled` audit row when it changed something. **Reconcile now** on the dashboard runs the same job on demand.
- Nothing deletes cases, documents or reports today, so counters only increase. If deletion is added it must decrement them (the nightly job would otherwise correct it the next night).

Migrating an existing single-tenant database

`alembic upgrade head` (as the owner role):

- creates `companies` / `company_usage_stats` and one company named `DEFAULT_COMPANY_NAME` ("Default Company");
- adds `company_id` everywhere and assigns **every** existing row to that company;
- turns existing `admin` accounts into `platform_admin` (no company); everyone else keeps their role in the default company;
- backfills the usage counters from the existing data;
- creates `fddt_app` / `fddt_platform` with the configured passwords, grants, and the RLS policies (tenant isolation for the app role, `platform_all` for the platform role).

Then `c9e3a1b5d7f4` (production hardening): removes `BYPASSRLS` from an existing `fddt_platform` (only if it has it), adds the `platform_all` policies, moves platform-access audit rows to platform-only, creates and seeds

`risk_rule_templates` , and enforces one row per check / page hash / assessment / signature match (collapsing duplicates earlier retries left: on the dev database 6 check rows, 1 page hash, 1 assessment).

The migrations need a table-owner role with **CREATEROLE**; no superuser.

`alembic downgrade` reverses it (keeping only the oldest company's rules/thresholds, since a single tenant can't hold several companies' copies).

Tests

Spec check	Where
Platform admin creates companies A and B and their users	<code>tests/test_multi_tenancy.py::tenants</code> fixture
A's reviewer_l2 sees A's cases, issuers, rules; 404 for every B resource by id (case detail/audit/file URL/signatures/reports/actions/issuer PATCH/rule history, <code>company_id=B</code> params)	<code>test_reviewer_l2_sees_own_company_and_gets_404_for_the_other</code> , <code>test_a_rule_history_of_another_company_is_a_404</code>
A's reviewer_l2 cannot reach Settings › Users	<code>test_reviewer_l2_cannot_reach_settings_users</code>
Platform admin opens B's case; a platform-only audit row is recorded (visible to platform admins, never to A or B)	<code>test_platform_admin_reads_any_case_and_each_access_is_audited</code>
Billing per company matches uploads; manual cross-check; drift corrected	<code>test_billing_dashboard_separates_companies_and_matches_uploads</code>
RLS alone blocks cross-tenant rows with no company filter in the SQL	<code>tests/test_qls_postgres.py</code> (needs <code>FDDT_RLS_TEST_DATABASE_URL</code>)
No BYPASSRLS anywhere; platform access via policies; platform-only audit rows invisible to the app role	<code>tests/test_qls_postgres.py</code> (<code>test_no_role_has_bypassrls</code> , <code>test_every_qls_table_has_the_platform_policy</code> , ...)
Company users never see platform-access rows (Audit History, filter, timeline)	<code>tests/test_multi_tenancy.py::test_company_users_never_see_platform_admin_access</code>
Every idempotent write (page hashes, check rows, usage counters) works under the least-privilege app role	<code>tests/test_qls_postgres.py::test_every_upsert_the_pipeline_uses_works_under_the_app_role</code>
New company gets the template; template edits don't change existing companies	<code>tests/test_rule_templates.py</code>
Tenant context never leaks through PgBouncer transaction pooling	<code>tests/test_pgbouncer_qls.py</code> (needs PgBouncer)
Pre-migration data intact and attributed to the default company	verified against the real dev database (see the change summary)