

FDDT Case Report

Fraud Document Detection Tool — system-assisted analysis for human review. See Section 6 for what this report is and is not.

1. Case Details

Case ID	CASE-3B02E9B2 3b02e9b2-2638-4e3c-8f3f-c4a591bd9621
Submitted by	testreviewer L!
Submission date	2026-10-07 06:44 UTC
Documents	1 Types: School document
Case type	School document
Status / decision	Submitted / Pending review
Report generated	2026-10-07 06:44 UTC by testreviewer L! (Reviewer L1)
Report ID	b1eac2cd-a7dc-49c5-a287-38dc2b42ddad

Preliminary report. Automated analysis had not finished when this report was generated. Still outstanding: Field validation of 'case1.pdf'; Issuer verification of 'case1.pdf'.

2. Executive Summary

Not yet scored. The risk engine has not produced an assessment for this case, so there is no overall tier.

Why: No risk tier yet — automated analysis is still in progress.

Decision status: Pending review.

Exceptions: 12 item(s) did not pass, 12 of them shown as highlighted regions on the source pages. Section 5 lists each one; Section 9 shows them on the documents.

This summary is deliberately short. The reasons behind the tier are in Section 3, the complete checklist is in Section 4.

3. Explainable Findings

The case has not been scored, so there are no triggered risk reasons yet.

4. All Checks

The completeness record: every check that ran on every document, including those that passed. Section 3 is the narrative; this is the checklist.

Document 1 — case1.pdf (School document · 4 passed, 3 flagged)

Check	Result	Details
Metadata forensics PDF metadata, edit history and incremental-save markers.	Pass	No anomaly detected by this check.
Font consistency Text set in a different font family from the text around it (text layer, or OCR font recognition on scans).	Flag	10 finding(s). Page 1 (scanned): '22,000' looks like a different font (grotesque sans) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned.
Error level analysis (ELA) Pixel-level recompression differences that can indicate edited regions.	Pass	No anomaly detected by this check.
Copy-move detection Regions of a page that were duplicated elsewhere on the same page.	Pass	No anomaly detected by this check.
Visual inconsistency / AI-generation review Model review of fonts, alignment, contrast, sharpness, lighting and signs of AI generation.	Flag	1 finding(s). Page 1 — Resolution / sharpness consistency: The pen signature and the round stamp near the bottom center are noticeably softer and less sharp compared to the printed text. This is a vision-model judgment, not pixel-level analysis — a probabilis...
Signature / stamp detection Whether a signature or stamp mark is present where this kind of document normally carries one.	Pass	2 signature/stamp region(s) located (presence only).
Duplicate detection Perceptual-hash comparison of each page against previously submitted pages.	Flag	1 finding(s). Page 1 matches page 1 of "sample_font20.pdf" in case CASE-37B33CBB, submitted by test1 Reviewer L2, and the two files are byte-for-byte identical (same SHA-256) — the exact same file re-uploaded.

Linked cases (reviewer note, not scored)

Other cases of this company that share a scanning device, a parent/guardian or a family/student ID with this one. Shared office scanners are common, so this is context for the reviewer only; it does not affect the risk score.

Case	Shared
CASE-11953D10	same scanning device (PPS - SFO2-A-MFP01, serial 00206BE9FB63); same parent/guardian ('Ahmed Salem Alkindi (father's name in HAMDA AHMED SALEM ALKINDI)' / 'Ahmed Salem Alkindi (father's name in HAMDA AHMED SALEM ALKINDI)'); same student ID (HSN1278)
CASE-11DD9B6A	same scanning device (PPS - SFO2-A-MFP01, serial 00206BE9FB63); same parent/guardian ('Ahmed Salem Alkindi (father's name in HAMDA AHMED SALEM ALKINDI)' / 'Mr. Ahmed Salem Abdulla AlKindi')
CASE-37B33CBB	same scanning device (PPS - SFO2-A-MFP01, serial 00206BE9FB63); same parent/guardian ('Ahmed Salem Alkindi (father's name in HAMDA AHMED SALEM ALKINDI)' / 'Ahmed Salem Alkindi (father's name in HAMDA AHMED SALEM ALKINDI)'); same student ID (HSN1278)
CASE-6CB8F7E1	same scanning device (PPS - SFO2-A-MFP01, serial 00206BE9FB63); same parent/guardian ('Ahmed Salem Alkindi (father's name in HAMDA AHMED SALEM ALKINDI)' / 'Ahmed Salem Alkindi (father's name in HAMDA AHMED SALEM ALKINDI)'); same student ID (HSN1278)
CASE-6CE43E0C	same scanning device (PPS - SFO2-A-MFP01, serial 00206BE9FB63); same parent/guardian ('Ahmed Salem Alkindi (father's name in HAMDA AHMED SALEM ALKINDI)' / 'Ahmed Salem Alkindi (father's name in HAMDA AHMED SALEM ALKINDI)'); same student ID (HSN1278)
CASE-7978A329	same scanning device (PPS - SFO2-A-MFP01, serial 00206BE9FB63); same parent/guardian ('Ahmed Salem Alkindi (father's name in HAMDA AHMED SALEM ALKINDI)' / 'Mr. Ahmed Salem Abdulla AlKindi'); same student ID (HSN1278)
CASE-92AA4164	same scanning device (PPS - SFO2-A-MFP01, serial 00206BE9FB63); same parent/guardian ('Ahmed Salem Alkindi (father's name in HAMDA AHMED SALEM ALKINDI)' / 'Ahmed Salem Alkindi (father's name in HAMDA AHMED SALEM ALKINDI)'); same student ID (HSN1278)

CASE-942C8493	same scanning device (PPS - SFO2-A-MFP01, serial 00206BE9FB63); same parent/guardian ('Ahmed Salem Alkindi (father's name in HAMDA AHMED SALEM ALKINDI)' / 'Ahmed Salem Alkindi (father's name in HAMDA AHMED SALEM ALKINDI)'); same student ID (HSN1278)
CASE-A7705426	same scanning device (PPS - SFO2-A-MFP01, serial 00206BE9FB63); same parent/guardian ('Ahmed Salem Alkindi (father's name in HAMDA AHMED SALEM ALKINDI)' / 'Mr. Ahmed Salem Abdulla AlKindi')
CASE-BDFC21C5	same scanning device (PPS - SFO2-A-MFP01, serial 00206BE9FB63); same parent/guardian ('Ahmed Salem Alkindi (father's name in HAMDA AHMED SALEM ALKINDI)' / 'Mr. Ahmed Salem Abdulla AlKindi')
CASE-CF609781	same scanning device (PPS - SFO2-A-MFP01, serial 00206BE9FB63); same parent/guardian ('Ahmed Salem Alkindi (father's name in HAMDA AHMED SALEM ALKINDI)' / 'Mr. Ahmed Salem Abdulla AlKindi')
CASE-DB43653A	same scanning device (PPS - SFO2-A-MFP01, serial 00206BE9FB63); same parent/guardian ('Ahmed Salem Alkindi (father's name in HAMDA AHMED SALEM ALKINDI)' / 'Ahmed Salem Alkindi (father's name in HAMDA AHMED SALEM ALKINDI)'); same student ID (HSN1278)
CASE-DE627FA2	same scanning device (PPS - SFO2-A-MFP01, serial 00206BE9FB63); same parent/guardian ('Ahmed Salem Alkindi (father's name in HAMDA AHMED SALEM ALKINDI)' / 'Mr. Ahmed Salem Abdulla AlKindi')

5. Exceptions

Everything that did **not** pass, drawn from Section 4 and written as concrete items (12 in total). A region ID (R1, R2, ...) points to the highlight that shows the exception on the source page in Section 9 — click it to jump there.

No.	Exception	Where	Highlight
1	Font consistency: text set in a different font family from the text around it (OCR font recognition on a scan) Page 1 (scanned): '22,000' looks like a different font (grotesque sans) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned.	Document 1 (case1.pdf) · page 1	R1
2	Font consistency: text set in a different font family from the text around it (OCR font recognition on a scan) Page 1 (scanned): 'MONTH' looks like a different font (grotesque sans) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned.	Document 1 (case1.pdf) · page 1	R2
3	Font consistency: text set in a different font family from the text around it (OCR font recognition on a scan) Page 1 (scanned): '4' looks like a different font (grotesque sans) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned.	Document 1 (case1.pdf) · page 1	R3
4	Font consistency: text set in a different font family from the text around it (OCR font recognition on a scan) Page 1 (scanned): '5,500' looks like a different font (grotesque sans) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned.	Document 1 (case1.pdf) · page 1	R4
5	Font consistency: text set in a different font family from the text around it (OCR font recognition on a scan) Page 1 (scanned): '16,500' looks like a different font (grotesque sans) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned.	Document 1 (case1.pdf) · page 1	R5
6	Font consistency: text set in a different font family from the text around it (OCR font recognition on a scan) Page 1 (scanned): '5,500' looks like a different font (grotesque sans) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned.	Document 1 (case1.pdf) · page 1	R6
7	Font consistency: text set in a different font family from the text around it (OCR font recognition on a scan) Page 1 (scanned): 'DUE' looks like a different font (monospace) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned. Only 1 word(s) on this page look like this font, so it is shown for review but not scored: OCR font recognition...	Document 1 (case1.pdf) · page 1	R7
8	Font consistency: text set in a different font family from the text around it (OCR font recognition on a scan) Page 1 (scanned): '16,600' looks like a different font (grotesque sans) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned.	Document 1 (case1.pdf) · page 1	R8
9	Font consistency: text set in a different font family from the text around it (OCR font recognition on a scan) Page 1 (scanned): '6,000' looks like a different font (grotesque sans) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned. Found on a second look at this page's numbers: OCR lists grotesque sans among the likely fonts for it (S...	Document 1 (case1.pdf) · page 1	R9
10	Font consistency: text set in a different font family from the text around it (OCR font recognition on a scan) Page 1 (scanned): '55,500' looks like a different font (grotesque sans) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned.	Document 1 (case1.pdf) · page 1	R10
11	Visual review: resolution / sharpness consistency inconsistency (model judgment) Page 1 — Resolution / sharpness consistency: The pen signature and the round stamp near the bottom center are noticeably softer and less sharp compared to the printed text. This is a vision-model judgment, not pixel-level analysis — a probabilistic signal, not a definitive finding. Not counted: it is about the detecte...	Document 1 (case1.pdf) · page 1	R11
12	Duplicate: page is a near-duplicate of an earlier submission Page 1 matches page 1 of "sample_font20.pdf" in case CASE-37B33CBB, submitted by test1Reviewer L2, and the two files are byte-for-byte identical (same SHA-256) — the exact same file re-uploaded.	Document 1 (case1.pdf) · page 1	R12

6. Limitations and Assumptions

Please read this section before relying on anything in this report.

Limitations of This Analysis

What this report is

This is **system-assisted analysis to support human review**. It is not a certified forensic or legal determination that any document is genuine, forged or fraudulent, and it must not be presented as one.

Three kinds of finding

Deterministic pixel and structure checks — error level analysis, copy-move detection and PDF metadata forensics — are computed by fixed algorithms over the file’s pixels and structure. Localized findings are drawn with **solid** boxes (ELA in red, copy-move in orange). They are signals, not proof: ELA in particular can false-positive on legitimately recompressed scans and can miss print-and-rescan tampering.

Deterministic rule-based field checks — field validation and cross-document consistency — compare extracted values by fixed rules and are drawn with **solid purple** boxes. They are exact about the values they compare, but only as reliable as the OCR and extraction that produced those values.

Model-judgment checks — the visual inconsistency review, the AI-generation assessment and the signature/stamp comparison — are **probabilistic assessments** made by a vision-language model, not deterministic proof. They can vary from run to run, and their boxes are only approximate; they are the **dashed** boxes captioned “approximate” throughout this report. Signature and stamp detection shows presence and placement only and is not an identity match; the AI-generation assessment is experimental, and its accuracy on scanned business documents has not been validated.

What the system can and cannot see

The system evaluates the **submitted document file itself**. It cannot verify facts external to the document — for example whether a real-world transaction actually occurred, whether the named issuer really produced the document, or (without a pre-existing reference signature) who a signatory is. Printed Arabic and English text are read; handwritten Arabic is not reliably read.

Absence of a finding

The absence of a flagged finding does not guarantee that a document is genuine. It means only that no supported check detected an anomaly.

Check coverage on this case

Not every check applies to every document or case. This is what actually ran here:

Check	Status	Coverage on this case
Metadata forensics	Ran	Ran on 1 of 1 document(s); 0 flagged.
Font consistency	Ran	Ran on 1 of 1 document(s); 1 flagged.
Error level analysis (ELA)	Ran	Ran on 1 of 1 document(s); 0 flagged.
Copy-move detection	Ran	Ran on 1 of 1 document(s); 0 flagged.
Visual inconsistency / AI-generation review	Ran	Ran on 1 of 1 document(s); 1 flagged.
Signature / stamp detection	Ran	Ran on 1 of 1 document(s); 0 flagged.
Duplicate detection	Ran	Ran on 1 of 1 document(s); 1 flagged.
Field validation	Did not run	Did not run on any of the 1 document(s).
Issuer verification	Did not run	Did not run on any of the 1 document(s).
Signature comparison (advisory)	Did not run	Did not run — no reference signature was set for this case.
Cross-document consistency	Not applicable	Not applicable — this case has a single document.

Risk score and rule versions

This case had not been scored when the report was generated, so no risk-rule versions apply.

Assumptions

What the analysis took to be true without independently verifying it. These are specific to what ran on this case.

- Field values were read by OCR followed by language-model extraction and are assumed to have been read correctly. Each field carries a confidence and an “uncertain” flag. This build records that flag but does not automatically route low-confidence documents to manual review, so it is listed here for the reviewer instead: no field was marked uncertain on this case.
- The SHA-256 chain of custody shows the stored file is unchanged since upload. It does not show that the uploaded file is the issuer’s original — a document altered before upload is judged on its content, like any other.
- Risk-rule weights, severities and tier thresholds are administrator-configured and assumed to be appropriately calibrated; the score is advisory and is not a probability of fraud.



7. Audit Trail

Chronological case-level events, read from the append-only audit log (times in UTC). This report does not write to or duplicate the log.

Time	Event	Actor	Detail
2026-10-07 06:44 UTC	Case created	testreviewer L! (Reviewer L1)	
2026-10-07 06:44 UTC	Document uploaded	testreviewer L! (Reviewer L1)	case1.pdf; SHA-256 910eb0b646a448aa...
2026-10-07 06:44 UTC	Metadata forensics completed	System	case1.pdf
2026-10-07 06:44 UTC	ELA and copy-move checks completed	System	case1.pdf
2026-10-07 06:44 UTC	Signature / stamp detection completed	System	case1.pdf
2026-10-07 06:44 UTC	Font consistency check completed	System	case1.pdf
2026-10-07 06:44 UTC	Text extraction, classification and field extraction completed	System	case1.pdf
2026-10-07 06:44 UTC	Visual inconsistency review completed	System	case1.pdf
2026-10-07 06:44 UTC	Duplicate detection completed	System	case1.pdf

8. Appendix

The detailed technical record, for someone auditing the analysis itself rather than a first-pass reviewer.

8.1 Extracted fields

Normalized values as stored at extraction time. “Located” is the page the value was matched to (used for the purple highlights); a blank means no position was recorded. Arabic values are shown right-to-left, numbers left-to-right.

Document 1 — case1.pdf

Field	Value	Confidence	Located
Issuer	HEADSTART NURSERY	95%	p.1
Reference number	10139/ 20220906	90%	p.1
Date	2022-09-06	95%	p.1
Amount	55,500.00 AED	90%	p.1
Subtotal	Not found		
Tax rate	Not found		
Tax amount	Not found		
Student name	HAMDA AHMED SALEM ALKINDI	95%	p.1
Student id	HSN1278	95%	p.1
Bank name	Abu Dhabi Commercial Bank (ADCB)	90%	p.1
Bank account number	10364749020001	90%	p.1
Iban	AE450030010364749020001	90%	p.1
Account title	HeadStart Nursery	90%	p.1
Address	PO Box 84158, Falaj Hazza, Al Ain, UAE	90%	p.1
Email	info@headstartnursery.ae	90%	p.1
Website	www.headstartnursery.ae	90%	p.1

8.2 Document hashes (chain of custody)

SHA-256 of each original as recorded when it was uploaded to immutable storage, re-checked against the stored file when this report was generated.

Document	SHA-256	Verification
case1.pdf 60,180 bytes	910eb0b646a448aa87ae22b0a1a4865f35c7a7b4956f971df204467f32f12af9	VERIFIED

8.3 Technical parameters used by each check

Page rendering (input to ELA, copy-move, duplicate, visual review)

Render resolution	200 DPI
Colour model	RGB, converted to BGR for OpenCV

Metadata forensics

Libraries	pikepdf / PyMuPDF (PDF structure, XMP, incremental updates)
Modified-after-created rule	flagged if ModDate is more than 1 second(s) after CreationDate
Editing-software names matched	photoshop, gimp, illustrator, affinity, coreldraw, paint.net, pixlr, canva, inkscape

Error level analysis (ELA)

JPEG recompression quality	75
Error gain (scale)	50 (sqrt error term; the reference tool's default)

Contrast stretch / LUT	not applied — the raw error magnitude is thresholded directly
Region threshold	page-relative percentile 99.7 with an absolute floor of 60/255
Region area accepted	0.2% – 60% of the page
Applies to	pages with an embedded raster image only (born-digital text pages are exempt)

Copy-move detection

Keypoint algorithm	BRISK (OpenCV), detector threshold 200
Response / matching / distance thresholds	90 / 20 / 15
Minimum cluster size	30 matching keypoint pairs
Minimum region area	2% of the page
Keypoint / match caps	4,000 / 4,000

Duplicate detection

Hash	perceptual hash (imagehash pHash), 64-bit, one per rendered page
Hamming-distance threshold	5 ($\leq$ this is flagged)
Compared against	every page hash of every previously submitted document

Visual inconsistency / AI-generation review

Model	gpt-4.1-mini
Independent runs per page	2
A finding is kept if	both runs report it, or one reports it with high confidence
Image sent to model	page render, long edge $\leq$ 1536px

Signature / stamp detection and comparison

Model	gpt-4.1-mini
Pages analysed per document	first 10
Output	presence and placement only; comparison is a qualitative verdict, never a numeric score

Field locations (purple highlights)

Method	extracted values matched back onto OCR word positions (Azure Document Intelligence Layout)
Repeated values	resolved by nearby keywords (“total”, “subtotal”, “VAT”, “date”...)
Accuracy	approximate; a field that cannot be matched is reported without a highlight

9. PDF Highlighted Regions

Every flagged page, rendered from the original with its findings drawn in. These are copies made for this report; the original files are untouched. Only pages with at least one highlight appear. Each highlight has a region ID that Section 5 refers back to.

Colour key

■	ELA (solid red)
■	Copy-move (solid orange)
■	Font mismatch in the text layer (solid fuchsia)
■	Field exception (solid purple)
■	Model judgment (dashed blue, approximate)
■	Page-level finding (not localized)

Region index

Region	Document	Page	What it shows
R1	Document 1 (case1.pdf)	1	Font mismatch — text set in a different font family (scan, estimated)
R2	Document 1 (case1.pdf)	1	Font mismatch — text set in a different font family (scan, estimated)
R3	Document 1 (case1.pdf)	1	Font mismatch — text set in a different font family (scan, estimated)
R4	Document 1 (case1.pdf)	1	Font mismatch — text set in a different font family (scan, estimated)
R5	Document 1 (case1.pdf)	1	Font mismatch — text set in a different font family (scan, estimated)
R6	Document 1 (case1.pdf)	1	Font mismatch — text set in a different font family (scan, estimated)
R7	Document 1 (case1.pdf)	1	Font mismatch — text set in a different font family (scan, estimated)
R8	Document 1 (case1.pdf)	1	Font mismatch — text set in a different font family (scan, estimated)
R9	Document 1 (case1.pdf)	1	Font mismatch — text set in a different font family (scan, estimated)
R10	Document 1 (case1.pdf)	1	Font mismatch — text set in a different font family (scan, estimated)
R11	Document 1 (case1.pdf)	1	Visual review — resolution / sharpness consistency (approximate)
R12	Document 1 (case1.pdf)	1	Duplicate detection — near-duplicate page

case1.pdf — page 1

Section 9 · Document 1 of 1 · regions R1, R2, R3, R4, R5, R6, R7, R8, R9, R10, R11, R12

Font mismatch in the text layer (solid fuchsia) Model judgment (dashed blue, approximate) Page-level finding (not localized)

حضانة هذ ستارت Headstart NURSERY

Sep 6, 2022 HEADSTART NURSERY

INVOICE Num: 10139/ 20220906

STUDENTS NAME: HAMDA AHMED SALEM ALKINDI

R7 · Font mismatch — text set in a different font family (scan, estimated)

R4 · Font mismatch — text set in a different font family (scan, estimated)

R3 · Font mismatch — text set in a different font family (scan, estimated)

R2 · Font mismatch — text set in a different font family (scan, estimated)

R1 · Font mismatch — text set in a different font family (scan, estimated)

FEE: 5,500 PER MONTH X 4 MONTH = 22,000

R6 · Font mismatch — text set in a different font family (scan, estimated)

R5 · Font mismatch — text set in a different font family (scan, estimated)

FEE: 5,500 PER MONTH X 3 MONTHS = 16,500

R9 · Font mismatch — text set in a different font family (scan, estimated)

R8 · Font mismatch — text set in a different font family (scan, estimated)

FEE: 5,000 PER MONTH X 3 MONTHS = 16,600

R10 · Font mismatch — text set in a different font family (scan, estimated)

55,500

DIRHAMS: TWENTY THOUSAND FIVE HUNDRED AED ONLY

R11 · Visual review — resolution / sharpness consistency (approximate)

Bank Name: Abu Dhabi Commercial Bank (ADCB)
Bank Account Number: 10364749020001
IBAN: AE450030010364749020001
Account Title: HeadStart Nursery

FOR HEADSTART NURSERY حضانة هذ ستارت

RAWAN TAQEZ HAZA AL-AIN - U.A.E.

PO Box 84158, Falaj Hazza, Al Ain, UAE | info@headstartnursery.ae | www.headstartnursery.ae

Scanned with CamScanner

R1. Font consistency

Page 1 (scanned): '22,000' looks like a different font (grotesque sans) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned.

R2. Font consistency

Page 1 (scanned): 'MONTH' looks like a different font (grotesque sans) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned.

R3. Font consistency

Page 1 (scanned): '4' looks like a different font (grotesque sans) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned.

R4. Font consistency

Page 1 (scanned): '5,500' looks like a different font (grotesque sans) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned.

R5. Font consistency

Page 1 (scanned): '16,500' looks like a different font (grotesque sans) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned.

R6. Font consistency

Page 1 (scanned): '5,500' looks like a different font (grotesque sans) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned.

R7. Font consistency

Page 1 (scanned): 'DUE' looks like a different font (monospace) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned. Only 1 word(s) on this page look like this font, so it is shown for review but not scored: OCR font recognition alone is not reliable on a word or two.

R8. Font consistency

Page 1 (scanned): '16,600' looks like a different font (grotesque sans) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned.

R9. Font consistency

Page 1 (scanned): '6,000' looks like a different font (grotesque sans) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned. Found on a second look at this page's numbers: OCR lists grotesque sans among the likely fonts for it (Segoe UI, Arial, Helvetica, sans-serif), as for the other amounts flagged here.

R10. Font consistency

Page 1 (scanned): '55,500' looks like a different font (grotesque sans) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned.

R11. Visual inconsistency / AI-generation review

Page 1 — Resolution / sharpness consistency: The pen signature and the round stamp near the bottom center are noticeably softer and less sharp compared to the printed text. This is a vision-model judgment, not pixel-level analysis — a probabilistic signal, not a definitive finding. Not counted: it is about the detected stamp (covers 100% of it) — pen strokes and rubber-stamp ink are naturally softer than printed text...

R12. Duplicate detection (page-level — not localized to a region)

Page 1 matches page 1 of "sample_font20.pdf" in case CASE-37B33CBB, submitted by test1Reviewer L2, and the two files are byte-for-byte identical (same SHA-256) — the exact same file re-uploaded.