

FDDT Case Report

Fraud Document Detection Tool — system-assisted analysis for human review. See Section 6 for what this report is and is not.

1. Case Details

Case ID	CASE-DB43653A db43653a-f104-457d-9bef-bd390a6d9f27
Submitted by	test1Reviewer L2
Submission date	2026-10-06 15:03 UTC
Documents	1 Types: School document
Case type	Commercial invoice
Status / decision	Pending manual review / Pending review
Report generated	2026-10-06 16:46 UTC by test1Reviewer L2 (Reviewer L2)
Report ID	d362f6b1-405c-4df5-a6d1-e3f6ee87e8c5

2. Executive Summary

HIGH RISK **Score 100 / 100** (computed 2026-10-06 15:04 UTC)

Why: Driven mainly by (40 points): 'case1.pdf': 1 page(s) are near-identical to previously submitted documents. Page 1 is a near-identical match (perceptual hash Hamming distance 0) to page 1 of "sample_font20.pdf" in case CASE-37B33C...

Decision status: Pending review.

Exceptions: 17 item(s) did not pass, 16 of them shown as highlighted regions on the source pages. Section 5 lists each one; Section 9 shows them on the documents.

This summary is deliberately short. The reasons behind the tier are in Section 3, the complete checklist is in Section 4.

3. Explainable Findings

The risk engine is a transparent, weighted rule set. Each row is one rule that fired, with the reason text stored at scoring time. Points are summed (raw 175), capped at 100, and mapped to a tier (medium from 30, high from 60) — here **100/100, high**.

Severity	Points	Check	Document / page	Reason
HIGH	40	Duplicate detection	case1.pdf · page 1	'case1.pdf': 1 page(s) are near-identical to previously submitted documents. Page 1 is a near-identical match (perceptual hash Hamming distance 0) to page 1 of "sample_font20.pdf" in case CASE-37B33CBB — a possible duplicate or near-duplicate resubmission. Highlights: R12
HIGH	25	Field validation	case1.pdf · page 1	'case1.pdf': The amount in words, 'TWENTY THOUSAND FIVE HUNDRED AED ONLY', is 20,500.00 (read from the words), but the total in figures is 55,500.00. Highlights: R13, R14, R15, R16
HIGH	25	Field validation	case1.pdf · page 1	'case1.pdf': 1 line item(s) do not add up: line 4 (FEE FOR TERM 3 (April-May-June)): $3 \times 6,000.00 = 18,000.00$, but the line total printed is 16,600.00. Highlights: R13, R14, R15, R16
MEDIUM	25	Font consistency	case1.pdf · page 1	'case1.pdf': on the scanned page, 8 word(s) look like a different font from the text around them (OCR font recognition — an estimate). Page 1 (scanned): '22,000' looks like a different font (grotesque sans) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned. Highlights: R1, R2, R3, R4, R5, R6, R7, R8, R9
MEDIUM	15	Field validation	case1.pdf · page 1	'case1.pdf': The line items add up to 33,100.00 (16,500.00 + 16,600.00), but the document states the total 55,500.00 — a difference of 22,400.00. Highlights: R13, R14, R15, R16
MEDIUM	15	Issuer verification	case1.pdf	Issuer 'HEADSTART NURSERY' on 'case1.pdf' did not match any known issuer in the registry.
MEDIUM	10	Metadata forensics	case1.pdf	'case1.pdf' embeds JavaScript or an open-action — unusual for a business document.
MEDIUM	10	Visual inconsistency / AI-generation review	case1.pdf · page 1	'case1.pdf': the vision-model review reports misaligned text. Page 1 — Text alignment: The text inside the box containing the student name and ID is slightly tilted and not aligned horizontally with the rest of the text on the page, around the middle top area (approx. $x=0.15$, $y=0.22$, $width=0.7$, $height=0.1$). This is a vision-model judgment, not pixel-level analysis — a probabilistic signal, not a definitive finding. Highlights: R10, R11
MEDIUM	10	Visual inconsistency / AI-generation review	case1.pdf · page 1	'case1.pdf': the vision-model review reports a region with different sharpness. Page 1 — Resolution / sharpness consistency: The stamp and signature near the bottom center of the page appear slightly more pixelated and lower resolution compared to the rest of the text, around ($x=0.3$, $y=0.75$, $width=0.4$, $height=0.1$). This is a vision-model judgment, not pixel-level analysis — a probabilistic signal, not a definitive finding. Highlights: R10, R11

4. All Checks

The completeness record: every check that ran on every document, including those that passed. Section 3 is the narrative; this is the checklist.

Document 1 — case1.pdf (School document · 3 passed, 6 flagged)

Check	Result	Details
Metadata forensics PDF metadata, edit history and incremental-save markers.	Flag	1 finding(s). This PDF contains an OpenAction and/or embedded JavaScript — unusual for a business document.
Font consistency Text set in a different font family from the text around it (text layer, or OCR font recognition on scans).	Flag	9 finding(s). Page 1 (scanned): '22,000' looks like a different font (grotesque sans) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned.
Error level analysis (ELA) Pixel-level recompression differences that can indicate edited regions.	Pass	No anomaly detected by this check.
Copy-move detection Regions of a page that were duplicated elsewhere on the same page.	Pass	No anomaly detected by this check.
Visual inconsistency / AI-generation review Model review of fonts, alignment, contrast, sharpness, lighting and signs of AI generation.	Flag	2 finding(s). Page 1 — Text alignment: The text inside the box containing the student name and ID is slightly tilted and not aligned horizontally with the rest of the text on the page, around the middle top area (approx. x=0.15, y=0.22, width=0.7, height=0.1)...
Signature / stamp detection Whether a signature or stamp mark is present where this kind of document normally carries one.	Pass	2 signature/stamp region(s) located (presence only).
Duplicate detection Perceptual-hash comparison of each page against previously submitted pages.	Flag	1 finding(s). Page 1 is a near-identical match (perceptual hash Hamming distance 0) to page 1 of "sample_font20.pdf" in case CASE-37B33CBB — a possible duplicate or near-duplicate resubmission.
Field validation Date, subtotal / tax / total and reference-number plausibility.	Flag	1 line item(s) do not add up: line 4 (FEE FOR TERM 3 (April-May-June)): 3 × 6,000.00 = 18,000.00, but the line total printed is 16,600.00. The amount in words, 'TWENTY THOUSAND FIVE HUNDRED AED ONLY', is 20,500.00 (read from the words), but the total in figur...
Issuer verification Issuer name matched against the issuer registry.	Flag	Issuer 'HEADSTART NURSERY' did not match any known issuer in the registry.

5. Exceptions

Everything that did **not** pass, drawn from Section 4 and written as concrete items (17 in total). A region ID (R1, R2, ...) points to the highlight that shows the exception on the source page in Section 9 — click it to jump there.

No.	Exception	Where	Highlight
1	Metadata forensics: This PDF contains an OpenAction and/or embedded JavaScript — unusual for a business document.	Document 1 (case1.pdf)	—
2	Font consistency: text set in a different font family from the text around it (OCR font recognition on a scan) Page 1 (scanned): '22,000' looks like a different font (grotesque sans) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned.	Document 1 (case1.pdf) · page 1	R1
3	Font consistency: text set in a different font family from the text around it (OCR font recognition on a scan) Page 1 (scanned): 'MONTH' looks like a different font (grotesque sans) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned.	Document 1 (case1.pdf) · page 1	R2
4	Font consistency: text set in a different font family from the text around it (OCR font recognition on a scan) Page 1 (scanned): '4' looks like a different font (grotesque sans) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned.	Document 1 (case1.pdf) · page 1	R3
5	Font consistency: text set in a different font family from the text around it (OCR font recognition on a scan) Page 1 (scanned): '5,500' looks like a different font (grotesque sans) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned.	Document 1 (case1.pdf) · page 1	R4
6	Font consistency: text set in a different font family from the text around it (OCR font recognition on a scan) Page 1 (scanned): '16,500' looks like a different font (grotesque sans) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned.	Document 1 (case1.pdf) · page 1	R5
7	Font consistency: text set in a different font family from the text around it (OCR font recognition on a scan) Page 1 (scanned): '5,500' looks like a different font (grotesque sans) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned.	Document 1 (case1.pdf) · page 1	R6
8	Font consistency: text set in a different font family from the text around it (OCR font recognition on a scan) Page 1 (scanned): 'DUE' looks like a different font (monospace) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned. Only 1 word(s) on this page look like this font, so it is shown for review but not scored: OCR font recognition...	Document 1 (case1.pdf) · page 1	R7
9	Font consistency: text set in a different font family from the text around it (OCR font recognition on a scan) Page 1 (scanned): '16,600' looks like a different font (grotesque sans) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned.	Document 1 (case1.pdf) · page 1	R8
10	Font consistency: text set in a different font family from the text around it (OCR font recognition on a scan) Page 1 (scanned): '55,500' looks like a different font (grotesque sans) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned.	Document 1 (case1.pdf) · page 1	R9
11	Visual review: text alignment inconsistency (model judgment) Page 1 — Text alignment: The text inside the box containing the student name and ID is slightly tilted and not aligned horizontally with the rest of the text on the page, around the middle top area (approx. x=0.15, y=0.22, width=0.7, height=0.1). This is a vision-model judgment, not pixel-level analysis — a probabilis...	Document 1 (case1.pdf) · page 1	R10
12	Visual review: resolution / sharpness consistency inconsistency (model judgment) Page 1 — Resolution / sharpness consistency: The stamp and signature near the bottom center of the page appear slightly more pixelated and lower resolution compared to the rest of the text, around (x=0.3, y=0.75, width=0.4, height=0.1). This is a vision-model judgment, not pixel-level analysis — a probabilistic signal...	Document 1 (case1.pdf) · page 1	R11
13	Duplicate: page is a near-duplicate of an earlier submission Page 1 is a near-identical match (perceptual hash Hamming distance 0) to page 1 of "sample_font20.pdf" in case CASE-37B33CBB — a possible duplicate or near-duplicate resubmission.	Document 1 (case1.pdf) · page 1	R12

14	Line item miscalculation: 1 line item(s) do not add up: line 4 (FEE FOR TERM 3 (April-May-June)): $3 \times 6,000.00 = 18,000.00$, but the line total printed is 16,600.00.	Document 1 (case1.pdf) · page 1	R13
15	Amount in words differs from figures: The amount in words, 'TWENTY THOUSAND FIVE HUNDRED AED ONLY', is 20,500.00 (read from the words), but the total in figures is 55,500.00.	Document 1 (case1.pdf) · page 1	R14, R15
16	Line items do not add up: The line items add up to 33,100.00 (16,500.00 + 16,600.00), but the document states the total 55,500.00 — a difference of 22,400.00.	Document 1 (case1.pdf) · page 1	R16
17	Issuer not in registry: Issuer 'HEADSTART NURSERY' did not match any known issuer in the registry.	Document 1 (case1.pdf)	—

2 exception(s) show “—” because they have no place on a page (for example a missing field, or a check that did not localize its finding).

6. Limitations and Assumptions

Please read this section before relying on anything in this report.

Limitations of This Analysis

What this report is

This is **system-assisted analysis to support human review**. It is not a certified forensic or legal determination that any document is genuine, forged or fraudulent, and it must not be presented as one.

Three kinds of finding

Deterministic pixel and structure checks — error level analysis, copy-move detection and PDF metadata forensics — are computed by fixed algorithms over the file’s pixels and structure. Localized findings are drawn with **solid** boxes (ELA in red, copy-move in orange). They are signals, not proof: ELA in particular can false-positive on legitimately recompressed scans and can miss print-and-rescan tampering.

Deterministic rule-based field checks — field validation and cross-document consistency — compare extracted values by fixed rules and are drawn with **solid purple** boxes. They are exact about the values they compare, but only as reliable as the OCR and extraction that produced those values.

Model-judgment checks — the visual inconsistency review, the AI-generation assessment and the signature/stamp comparison — are **probabilistic assessments** made by a vision-language model, not deterministic proof. They can vary from run to run, and their boxes are only approximate; they are the **dashed** boxes captioned “approximate” throughout this report. Signature and stamp detection shows presence and placement only and is not an identity match; the AI-generation assessment is experimental, and its accuracy on scanned business documents has not been validated.

What the system can and cannot see

The system evaluates the **submitted document file itself**. It cannot verify facts external to the document — for example whether a real-world transaction actually occurred, whether the named issuer really produced the document, or (without a pre-existing reference signature) who a signatory is. Printed Arabic and English text are read; handwritten Arabic is not reliably read.

Absence of a finding

The absence of a flagged finding does not guarantee that a document is genuine. It means only that no supported check detected an anomaly.

Check coverage on this case

Not every check applies to every document or case. This is what actually ran here:

Check	Status	Coverage on this case
Metadata forensics	Ran	Ran on 1 of 1 document(s); 1 flagged.
Font consistency	Ran	Ran on 1 of 1 document(s); 1 flagged.
Error level analysis (ELA)	Ran	Ran on 1 of 1 document(s); 0 flagged.
Copy-move detection	Ran	Ran on 1 of 1 document(s); 0 flagged.
Visual inconsistency / AI-generation review	Ran	Ran on 1 of 1 document(s); 1 flagged.
Signature / stamp detection	Ran	Ran on 1 of 1 document(s); 0 flagged.
Duplicate detection	Ran	Ran on 1 of 1 document(s); 1 flagged.
Field validation	Ran	Ran on 1 of 1 document(s); 1 flagged.
Issuer verification	Ran	Ran on 1 of 1 document(s); 1 flagged.
Signature comparison (advisory)	Did not run	Did not run — no reference signature was set for this case.
Cross-document consistency	Not applicable	Not applicable — this case has a single document.

Risk score and rule versions

The risk score in this report (**100/100, high**) was computed at 2026-10-06 15:04 UTC (assessment 7eb4e90e-320f-4d1d-b7c6-bfcdefe437f9) from the risk-rule versions that fired: duplicate.cross_case_match v1 (weight 40); field.amount_in_words_mismatch v1 (weight 25); field.line_item_arithmetic_mismatch v1 (weight 25);

field.subtotal_line_item_mismatch v1 (weight 15); font.inconsistency_scanned v1 (weight 25); issuer.not_in_registry v1 (weight 15); metadata.javascript_or_openaction v1 (weight 10); visual.alignment_inconsistency v1 (weight 10); visual.sharpness_inconsistency v1 (weight 10). Tier thresholds at that time: medium from 30, high from 60.

Rule weights are versioned and are never re-applied retroactively: if an administrator later edits a rule, this case is not re-scored, and **this report remains an accurate historical record** of what was known and weighted when it was generated. A report generated later may legitimately differ.

Assumptions

What the analysis took to be true without independently verifying it. These are specific to what ran on this case.

- The issuer registry (0 active entries) was assumed to be current and accurate as of 2026-10-06 16:46 UTC. An issuer that is absent from it is reported as “not in registry”, which is a prompt to verify the issuer — not a finding that the issuer is fraudulent.
- Field values were read by OCR followed by language-model extraction and are assumed to have been read correctly. Each field carries a confidence and an “uncertain” flag. This build records that flag but does not automatically route low-confidence documents to manual review, so it is listed here for the reviewer instead: no field was marked uncertain on this case.
- Purple field highlights were located by matching extracted values back onto OCR word positions; treat their position as approximate. The comparison behind each one uses the stored, normalized values, not the highlight.
- The SHA-256 chain of custody shows the stored file is unchanged since upload. It does not show that the uploaded file is the issuer’s original — a document altered before upload is judged on its content, like any other.
- Risk-rule weights, severities and tier thresholds are administrator-configured and assumed to be appropriately calibrated; the score is advisory and is not a probability of fraud.

7. Audit Trail

Chronological case-level events, read from the append-only audit log (times in UTC). This report does not write to or duplicate the log.

Time	Event	Actor	Detail
2026-10-06 15:03 UTC	Case created	test1Reviewer L2 (Reviewer L2)	
2026-10-06 15:03 UTC	Document uploaded	test1Reviewer L2 (Reviewer L2)	case1.pdf; SHA-256 910eb0b646a448aa...
2026-10-06 15:03 UTC	Metadata forensics completed	System	case1.pdf
2026-10-06 15:03 UTC	ELA and copy-move checks completed	System	case1.pdf
2026-10-06 15:04 UTC	Signature / stamp detection completed	System	case1.pdf
2026-10-06 15:04 UTC	Duplicate detection completed	System	case1.pdf
2026-10-06 15:04 UTC	Font consistency check completed	System	case1.pdf
2026-10-06 15:04 UTC	Text extraction, classification and field extraction completed	System	case1.pdf
2026-10-06 15:04 UTC	Visual inconsistency review completed	System	case1.pdf
2026-10-06 15:04 UTC	Field validation and issuer verification completed	System	case1.pdf
2026-10-06 15:04 UTC	Risk score computed	System	tier high, score 100
2026-10-06 15:04 UTC	Case status changed	System	submitted → pending_manual_review; automated checks complete

8. Appendix

The detailed technical record, for someone auditing the analysis itself rather than a first-pass reviewer.

8.1 Extracted fields

Normalized values as stored at extraction time. “Located” is the page the value was matched to (used for the purple highlights); a blank means no position was recorded. Arabic values are shown right-to-left, numbers left-to-right.

Document 1 — case1.pdf

Field	Value	Confidence	Located
Issuer	HEADSTART NURSERY	95%	p.1
Reference number	10139/ 20220906	90%	p.1
Date	2022-09-06	95%	p.1
Amount	55,500.00 AED	95%	p.1
Subtotal	Not found		
Tax rate	Not found		
Tax amount	Not found		
Student name	HAMDA AHMED SALEM ALKINDI	95%	p.1
Student id	HSN1278	95%	p.1
Bank name	Abu Dhabi Commercial Bank (ADCB)	90%	p.1
Bank account number	10364749020001	90%	p.1
Iban	AE450030010364749020001	90%	p.1
Account title	HeadStart Nursery	90%	p.1
Address	PO Box 84158, Falaj Hazza, Al Ain, UAE	90%	p.1
Email	info@headstartnursery.ae	90%	p.1
Website	www.headstartnursery.ae	90%	p.1

8.2 Document hashes (chain of custody)

SHA-256 of each original as recorded when it was uploaded to immutable storage, re-checked against the stored file when this report was generated.

Document	SHA-256	Verification
case1.pdf 60,180 bytes	910eb0b646a448aa87ae22b0a1a4865f35c7a7b4956f971df204467f32f12af9	VERIFIED

8.3 Technical parameters used by each check

Page rendering (input to ELA, copy-move, duplicate, visual review)

Render resolution	200 DPI
Colour model	RGB, converted to BGR for OpenCV

Metadata forensics

Libraries	pikepdf / PyMuPDF (PDF structure, XMP, incremental updates)
Modified-after-created rule	flagged if ModDate is more than 1 second(s) after CreationDate
Editing-software names matched	photoshop, gimp, illustrator, affinity, coreldraw, paint.net, pixlr, canva, inkscape

Error level analysis (ELA)

JPEG recompression quality	75
Error gain (scale)	50 (sqrt error term; the reference tool's default)

Contrast stretch / LUT	not applied — the raw error magnitude is thresholded directly
Region threshold	page-relative percentile 99.7 with an absolute floor of 60/255
Region area accepted	0.2% – 60% of the page
Applies to	pages with an embedded raster image only (born-digital text pages are exempt)

Copy-move detection

Keypoint algorithm	BRISK (OpenCV), detector threshold 200
Response / matching / distance thresholds	90 / 20 / 15
Minimum cluster size	30 matching keypoint pairs
Minimum region area	2% of the page
Keypoint / match caps	4,000 / 4,000

Duplicate detection

Hash	perceptual hash (imagehash pHash), 64-bit, one per rendered page
Hamming-distance threshold	5 ($\leq$ this is flagged)
Compared against	every page hash of every previously submitted document

Visual inconsistency / AI-generation review

Model	gpt-4.1-mini
Independent runs per page	2
A finding is kept if	both runs report it, or one reports it with high confidence
Image sent to model	page render, long edge $\leq$ 1536px

Signature / stamp detection and comparison

Model	gpt-4.1-mini
Pages analysed per document	first 10
Output	presence and placement only; comparison is a qualitative verdict, never a numeric score

Field validation

Sub-checks	date not in the future; subtotal vs line items; total vs subtotal + tax; reference-number format
Amount tolerance	1% of the amount, minimum 0.01

Issuer verification

Matching	rapidfuzz weighted ratio against the active issuer registry, with a model-judged fallback for cross-script names
Similarity threshold	85

Field locations (purple highlights)

Method	extracted values matched back onto OCR word positions (Azure Document Intelligence Layout)
Repeated values	resolved by nearby keywords (“total”, “subtotal”, “VAT”, “date”...)
Accuracy	approximate; a field that cannot be matched is reported without a highlight

9. PDF Highlighted Regions

Every flagged page, rendered from the original with its findings drawn in. These are copies made for this report; the original files are untouched. Only pages with at least one highlight appear. Each highlight has a region ID that Section 5 refers back to.

Colour key


■	ELA (solid red)
■	Copy-move (solid orange)
■	Font mismatch in the text layer (solid fuchsia)
■	Field exception (solid purple)
■	Model judgment (dashed blue, approximate)
■	Page-level finding (not localized)

Region index

Region	Document	Page	What it shows
R1	Document 1 (case1.pdf)	1	Font mismatch — text set in a different font family (scan, estimated)
R2	Document 1 (case1.pdf)	1	Font mismatch — text set in a different font family (scan, estimated)
R3	Document 1 (case1.pdf)	1	Font mismatch — text set in a different font family (scan, estimated)
R4	Document 1 (case1.pdf)	1	Font mismatch — text set in a different font family (scan, estimated)
R5	Document 1 (case1.pdf)	1	Font mismatch — text set in a different font family (scan, estimated)
R6	Document 1 (case1.pdf)	1	Font mismatch — text set in a different font family (scan, estimated)
R7	Document 1 (case1.pdf)	1	Font mismatch — text set in a different font family (scan, estimated)
R8	Document 1 (case1.pdf)	1	Font mismatch — text set in a different font family (scan, estimated)
R9	Document 1 (case1.pdf)	1	Font mismatch — text set in a different font family (scan, estimated)
R10	Document 1 (case1.pdf)	1	Visual review — text alignment (approximate)
R11	Document 1 (case1.pdf)	1	Visual review — resolution / sharpness consistency (approximate)
R12	Document 1 (case1.pdf)	1	Duplicate detection — near-duplicate page
R13	Document 1 (case1.pdf)	1	Field exception: Line 4 total (16,600.00 — 3 × 6,000.00 = 18,000.00)
R14	Document 1 (case1.pdf)	1	Field exception: Total (55,500.00 vs 20,500.00 in words)
R15	Document 1 (case1.pdf)	1	Field exception: Amount in words (20,500.00 vs total 55,500.00)
R16	Document 1 (case1.pdf)	1	Field exception: Total amount (55,500.00 AED vs line items 33,100.00)

case1.pdf — page 1

Section 9 · Document 1 of 1 · regions R1, R2, R3, R4, R5, R6, R7, R8, R9, R10, R11, R12, R13, R14, R15, R16

 Font mismatch in the text layer (solid fuchsia) Field exception (solid purple) Model judgment (dashed blue, approximate) Page-level finding

The image shows a scanned receipt from Headstart Nursery. The receipt is in Arabic and English. It includes the nursery's logo, date (Sep 6, 2022), and a list of items with prices. There are several annotations in purple and blue boxes over the receipt, indicating font mismatches and field exceptions. The receipt also includes a bank account number and a signature.

Headstart NURSERY

Sep 6, 2022

R16 · Field exception: Total amount (55,500.00 AED vs line items 33,100.00)

R10 · Visual review — text alignment (approximate)

R14 · Field exception: Total (55,500.00 vs 20,500.00 in words)

R7 · Font mismatch — text set in a different font family (scan, estimated)

R4 · Font mismatch — text set in a different font family (scan, estimated)

R3 · Font mismatch — text set in a different font family (scan, estimated)

R2 · Font mismatch — text set in a different font family (scan, estimated)

R1 · Font mismatch — text set in a different font family (scan, estimated)

FEE: 5,500 PER MONTH X 4 MONTH = 22,000

R6 · Font mismatch — text set in a different font family (scan, estimated)

R5 · Font mismatch — text set in a different font family (scan, estimated)

FEE: 5,500 PER MONTH X 3 MONTHS = 16,500

R13 · Field exception: Line 4 total (16,600.00 — 3 × 6,000.00 = 18,000.00)

R8 · Font mismatch — text set in a different font family (scan, estimated)

FEE: 6,000 PER MONTH X 3 MONTHS = 18,000

R9 · Font mismatch — text set in a different font family (scan, estimated)

R15 · Field exception: Amount in words (20,500.00 vs total 55,500.00)

DIRHAMS: TWENTY THOUSAND FIVE HUNDRED AED ONLY

R11 · Visual review — resolution / sharpness consistency (approximate)

Bank Name: Abu Dhabi Commercial Bank (ADCB)

Bank Account Number: 10364749020001

IBAN: AE450030010364749020001

Account Title: HeadStart Nursery

FOR HEADSTART NURSERY

RAWAN TAQEZ

PO Box 84158, Falaj Hazza, Al Ain, UAE | info@headstartnursery.ae | www.headstartnursery.ae

Scanned with CamScanner

R1. Font consistency

Page 1 (scanned): '22,000' looks like a different font (grotesque sans) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned.

R2. Font consistency

Page 1 (scanned): 'MONTH' looks like a different font (grotesque sans) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned.

R3. Font consistency

Page 1 (scanned): '4' looks like a different font (grotesque sans) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned.

R4. Font consistency

Page 1 (scanned): '5,500' looks like a different font (grotesque sans) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned.

R5. Font consistency

Page 1 (scanned): '16,500' looks like a different font (grotesque sans) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned.

R6. Font consistency

Page 1 (scanned): '5,500' looks like a different font (grotesque sans) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned.

R7. Font consistency

Page 1 (scanned): 'DUE' looks like a different font (monospace) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned. Only 1 word(s) on this page look like this font, so it is shown for review but not scored: OCR font recognition alone is not reliable on a word or two.

R8. Font consistency

Page 1 (scanned): '16,600' looks like a different font (grotesque sans) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned.

R9. Font consistency

Page 1 (scanned): '55,500' looks like a different font (grotesque sans) from the text around it (humanist sans), per OCR font recognition — it may have been typed in or pasted before the page was printed or scanned.

R10. Visual inconsistency / AI-generation review

Page 1 — Text alignment: The text inside the box containing the student name and ID is slightly tilted and not aligned horizontally with the rest of the text on the page, around the middle top area (approx. x=0.15, y=0.22, width=0.7, height=0.1). This is a vision-model judgment, not pixel-level analysis — a probabilistic signal, not a definitive finding.

R11. Visual inconsistency / AI-generation review

Page 1 — Resolution / sharpness consistency: The stamp and signature near the bottom center of the page appear slightly more pixelated and lower resolution compared to the rest of the text, around (x=0.3, y=0.75, width=0.4, height=0.1). This is a vision-model judgment, not pixel-level analysis — a probabilistic signal, not a definitive finding.

R12. Duplicate detection (page-level — not localized to a region)

Page 1 is a near-identical match (perceptual hash Hamming distance 0) to page 1 of "sample_font20.pdf" in case CASE-37B33CBB — a possible duplicate or near-duplicate resubmission.

R13. Field validation

1 line item(s) do not add up: line 4 (FEE FOR TERM 3 (April-May-June)): 3 × 6,000.00 = 18,000.00, but the line total printed is 16,600.00. (Line 4 total: 16,600.00)

R14. Field validation

The amount in words, 'TWENTY THOUSAND FIVE HUNDRED AED ONLY', is 20,500.00 (read from the words), but the total in figures is 55,500.00. (Total amount: 55,500.00 AED)

R15. Field validation

The amount in words, 'TWENTY THOUSAND FIVE HUNDRED AED ONLY', is 20,500.00 (read from the words), but the total in figures is 55,500.00. (Amount in words: TWENTY THOUSAND FIVE HUNDRED AED ONLY)

R16. Field validation

The line items add up to 33,100.00 (16,500.00 + 16,600.00), but the document states the total 55,500.00 — a difference of 22,400.00. (Total amount: 55,500.00 AED)